

REPORT

Abhay Vivek Kulkarni (ak6277)

Project 2: Bitonic Sort

Program Workflow:

1) Get user input for size of raw data.

--size-N : Generates an array of size N with random numbers
--print-X : (optional) Prints a part of the sorted array of size X for easy readability.
If not specified; does not print the result

Global parameters: (hardcoded)

MASTER = Specifies the base thread

SEED = Seed for the random value generation

BOUND = range of random values (0 to BOUND)

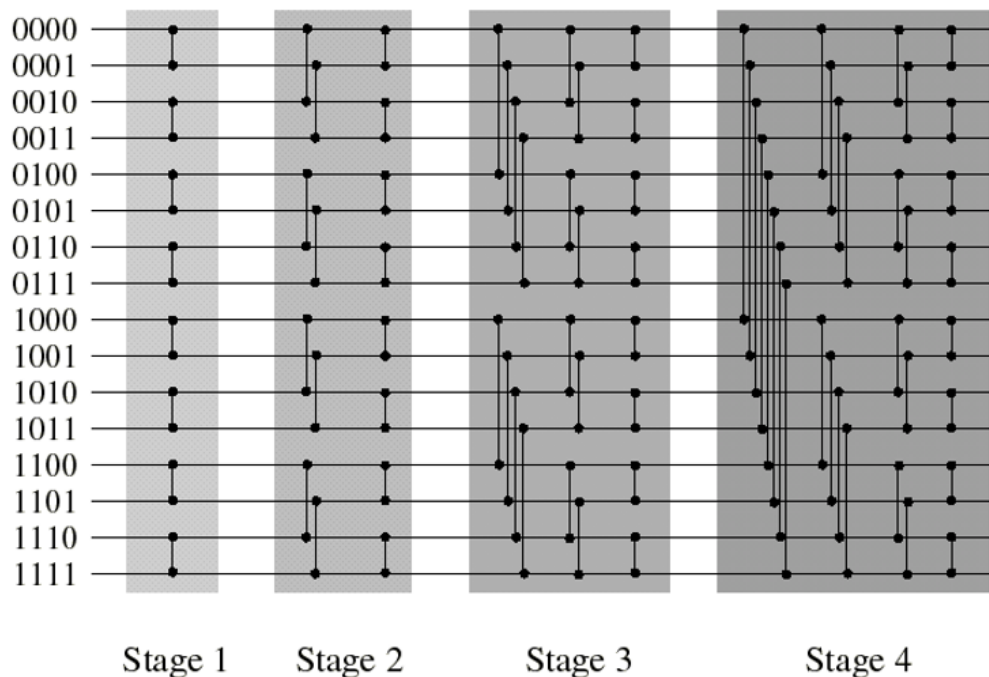
2) MPI.Init()

All processors get (N/P) elements

Wait for all processors to get their raw data

Stages = $\log_2 P$

Processors



3) For each stage:

Perform comparisons:

if (stage is even) and (processor's index = 0) OR

if (stage is odd) and (processor's index = 1)

Perform a min comparison

else

Perform a max comparison

```
for (i = 0; i < stages; i++) {  
    for (j = i; j >= 0; j--) {  
        if (((process_rank >> (i + 1)) % 2 == 0 && (process_rank >> j) % 2 == 0) ||  
            ((process_rank >> (i + 1)) % 2 != 0 && (process_rank >> j) % 2 != 0)) {  
            minCompare(j);  
        } else {  
            maxCompare(j);  
        }  
    }  
}
```

This transforms the Raw data to Bitonic by stage 3 (in this case)

The last stage converts Bitonic sequence to sorted sequence.

4) Min-Max Comparisons:

process rank P1 compares its value with (P1 XOR (left bit shift))

swaps values based on the type of comparison

<eg> For 16 processors (Assigned 0000 to 1111)

Stage 1: 0000 - 0001 0100 - 0101

Stage 2: 0000 - 0010 0100 - 0110

Stage 3: 0000 - 0100 0100 - 0000

Stage 4: 0000 - 1000 0100 - 1100

Observations and Issues:

1) No significant changes can be seen in execution time for array sizes < 4M
as spawning threads in Java has an overhead which counteracts the performance gain achieved by multi-processing.

2) For <1024 sizes, regular in-built sort is extremely efficient.

3) Significant changes can be observed at array size of ~ 100 M.
Sequential is 7.5 times slower than parallel 8-core processing.

4) With input size > 268 M (16384^2)

The final master array can't accommodate that many values.

Issues with program heap/ int overflows

Assigning (long) values would add to the extra memory overhead.

Result graph:

